

Discrete Classification

1 Some preliminaries

There are a few useful matrix properties we will use in this lecture. These include the following:

$$\begin{aligned}\frac{d}{d\mathbf{x}} \mathbf{x}^T \mathbf{A} \mathbf{x} &= (\mathbf{A} + \mathbf{A}^T) \mathbf{x} \\ \frac{d}{d\mathbf{X}} \text{Tr}(\mathbf{X}^{-1} \mathbf{A}) &= -\mathbf{X}^{-T} \mathbf{A}^T \mathbf{X}^{-T} \\ \frac{d}{d\mathbf{X}} \log |\mathbf{X}| &= \mathbf{X}^{-T}\end{aligned}$$

We will also use the fact that for matrices A, B, C, D , we have that

$$\text{Tr}(ABCD \dots) = \text{Tr}(BCD \dots A) = \text{Tr}(CD \dots AB)$$

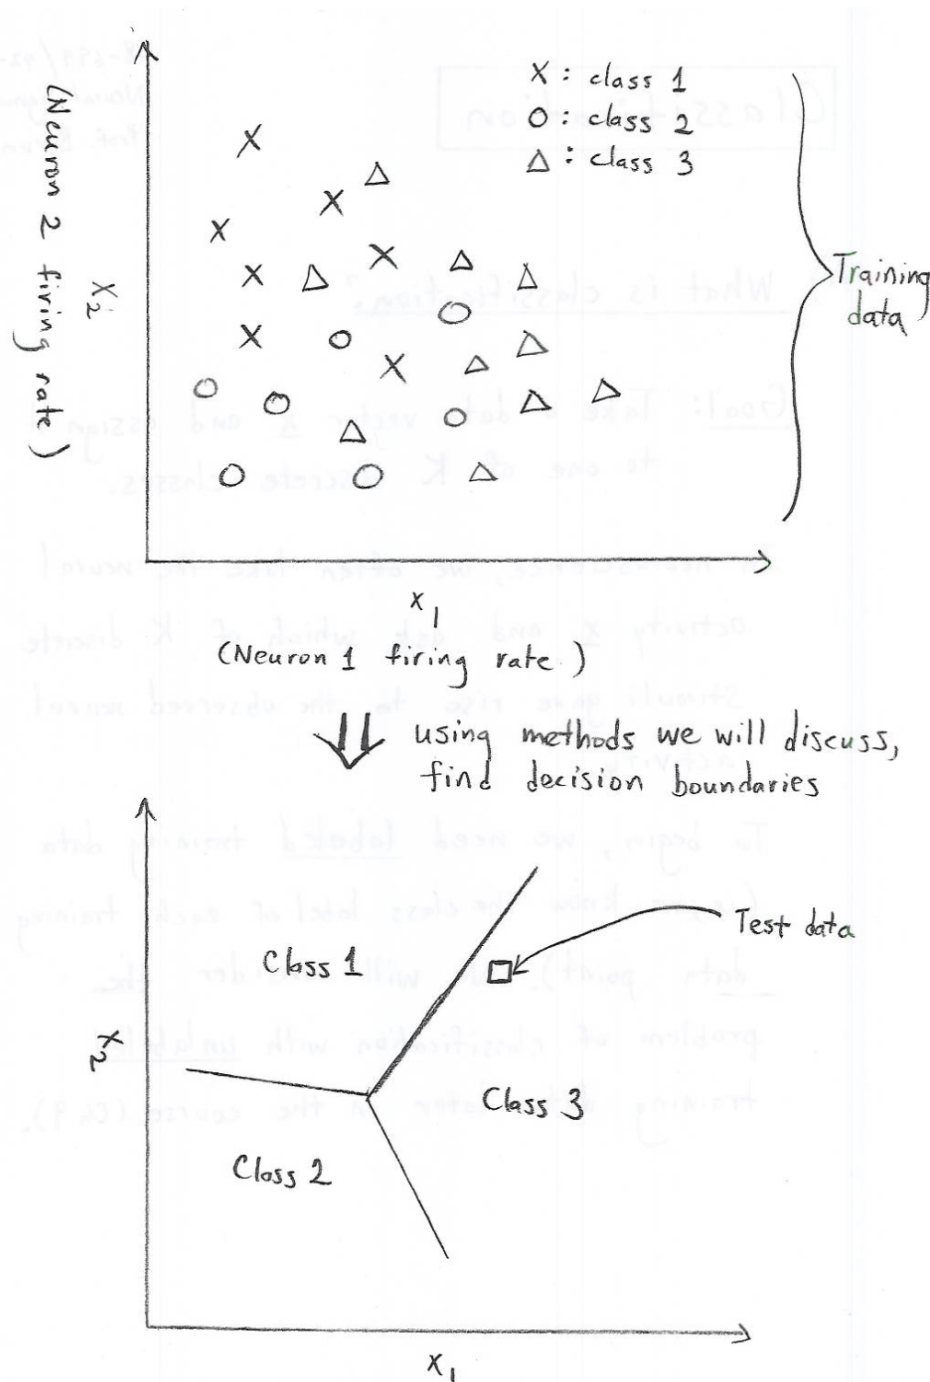
In general, a good reference is available at: <http://www.ee.ic.ac.uk/hp/staff/dmb/matrix/intro.html> or simply google “matrix reference manual.”

2 What is classification?

The goal of classification is to take a data vector \mathbf{x} and assign it to one of K discrete classes.

In neuroscience, we often take the neural activity \mathbf{x} and ask which of k discrete stimuli gave rise to the observed neural activity.

To begin, we need *labeled* training data (i.e., we know the class label of each training data point). We will consider the problem of classification with *unlabeled* training data later on in the course.



3 Classifying using generative models

There are two phases during classification: training and testing. (In machine learning formally, there are usually training, validation, and testing sets. For now, we'll

simplify the problem by just considering training and testing sets.)

In the **training phase**, we fit class-conditional densities $P(\mathbf{x}|C_k)$ and class priors $P(C_k)$ to training data (for $k = 1, \dots, K$).

In the **testing phase**, we:

- Compute $P(C_k|\mathbf{x})$, where \mathbf{x} is the test data, using Bayes rule, i.e.,

$$\begin{aligned} P(C_k|\mathbf{x}) &= \frac{P(\mathbf{x}|C_k)P(C_k)}{P(\mathbf{x})} \\ &= \frac{P(\mathbf{x}|C_k)P(C_k)}{\sum_{j=1}^K P(\mathbf{x}|C_j)P(C_j)} \end{aligned}$$

- Decode the class $\hat{k} = \arg \max_k P(C_k|\mathbf{x})$ to test data \mathbf{x} .

4 Generative models

The probabilities $P(\mathbf{x}|C_k)$ and $P(C_k)$ define a “probabilistic generative model.” This means that we can generate synthetic data from the model.

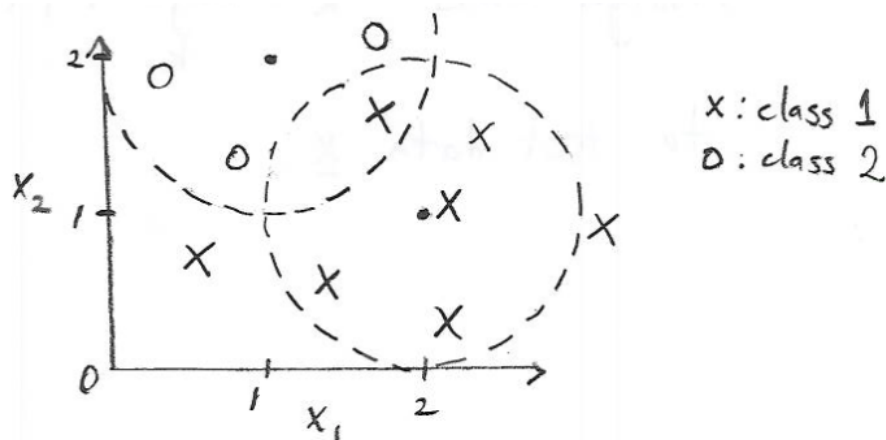
For example, say there are two classes and $\mathbf{x} \in \mathbb{R}^2$, with

$$\begin{aligned} P(C_1) &= 0.7 \\ P(C_2) &= 0.3 \\ P(\mathbf{x}|C_1) &= \mathcal{N}\left(\begin{bmatrix} 2 \\ 1 \end{bmatrix}, \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}\right) \\ P(\mathbf{x}|C_2) &= \mathcal{N}\left(\begin{bmatrix} 1 \\ 2 \end{bmatrix}, \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}\right) \end{aligned}$$

To generate one synthetic data vector \mathbf{x} , first flip a biased coin with probability 0.7 of coming up heads, and,

- If heads, draw from the Gaussian $P(\mathbf{x}|C_1)$.
- If tails, draw from the Gaussian $P(\mathbf{x}|C_2)$.

The philosophy of generative models is as follows. If we generate synthetic data from the model, and it looks a lot like the real data we’re trying to model, then we have a good model for our real data. We can then use the generative model to make optimal inferences, decisions, etc.



5 Training phase: maximum-likelihood and parameter estimation

One common technique to train models using machine learning is to maximize the likelihood of the data with respect to the model parameters. The idea is that, given a model, the best parameters are those under which the observed data was most probable. We'll go through a concrete example of this.

Example: Two classes with Gaussian class-conditional density with shared covariance.

Training data: $\{\mathbf{x}_n, t_n\}$ for $n = 1, \dots, N$.

- $t_n = 1$ denotes class C_1 .
- $t_n = 0$ denotes class C_2 .

We also set,

$$\begin{aligned} \Pr(t_n = 1) &= P(C_1) \\ &= \pi \\ \Pr(t_n = 0) &= P(C_2) \\ &= 1 - \pi \end{aligned}$$

Let's consider a data point, $\mathbf{x}_n \in \mathbb{R}^D$. The probability of having observed this data

point and it coming from C_1 is

$$\begin{aligned} P(\mathbf{x}_n, C_1) &= P(\mathbf{x}_n|C_1)P(C_1) \\ &= \mathcal{N}(\mathbf{x}_n|\mu_1, \Sigma)\pi \end{aligned}$$

Similarly, the probability of having observed this data point and it coming from C_2 is

$$\begin{aligned} P(\mathbf{x}_n, C_2) &= P(\mathbf{x}_n|C_2)P(C_2) \\ &= \mathcal{N}(\mathbf{x}_n|\mu_2, \Sigma)(1 - \pi) \end{aligned}$$

In the training set, we observe N data points together, as well as their classes. Therefore, the *likelihood* of this data under our model is given by:

$$\begin{aligned} \mathcal{L} &= P(\{\mathbf{x}_n, t_n\}|\pi, \mu_1, \mu_2, \Sigma) \\ &= \prod_{i=1}^N (\mathcal{N}(\mathbf{x}_n|\mu_1, \Sigma)\pi)^{t_n} (\mathcal{N}(\mathbf{x}_n|\mu_2, \Sigma)(1 - \pi))^{1-t_n} \end{aligned}$$

Products and exponentials are often inconvenient to deal with, and so we deal with this by taking the log of the likelihood. Because $\log(\cdot)$ is a monotonically increasing function of its argument, maximizing $\log(\mathcal{L})$ is equivalent to maximizing \mathcal{L} .

Concretely,

$$\begin{aligned} \log \mathcal{L} &= \sum_{n=1}^N (t_n \log \mathcal{N}(\mathbf{x}_n|\mu_1, \Sigma) + t_n \log \pi + \dots \\ &\quad + (1 - t_n) \log \mathcal{N}(\mathbf{x}_n|\mu_2, \Sigma) + (1 - t_n) \log(1 - \pi)) \end{aligned}$$

where

$$\log \mathcal{N}(\mathbf{x}_n|\mu_k, \Sigma) = -\frac{1}{2}(\mathbf{x}_n - \mu_k)^T \Sigma^{-1}(\mathbf{x}_n - \mu_k) - \frac{1}{2} \log |\Sigma| - \frac{D}{2} \log(2\pi)$$

Next, we want to find the optimal parameters $\pi, \mu_1, \mu_2, \Sigma$ are.

Now, we find each parameter via optimization.